

compboost

Fast and Flexible Component-Wise Boosting

Daniel Schalk

November 29, 2018



Use-Case

- We own a small booth at the Christmas market that sells mulled wine.
- As we are very interested in our customers' health, we only sell to customers who we expect to drink less than 15 liters per season.
- To estimate how much a customer drinks, we have collected data from 200 customers in recent years.
- These data include mulled wine consumption (in liter and cups), age, sex, country of origin, weight, body size, and 200 characteristics (that have absolutely no influence).

mw_consumption	mw_consumption_cups	gender	country	age	weight	height	noise1
12.6	42	f	Seychelles	21	119.25	157.9	0.1680
2.1	7	f	Poland	57	67.72	157.5	0.8075
12.6	42	m	Seychelles	25	65.54	181.6	0.3849
9.3	31	f	Germany	68	84.81	195.9	0.3277
4.8	16	m	Ireland	39	85.30	163.9	0.6021
9.3	31	f	Germany	19	102.75	173.8	0.6044
5.1	17	f	Poland	66	98.54	199.7	0.1246
7.5	25	f	Seychelles	62	108.21	198.1	0.2946
15.0	50	m	Seychelles	22	62.08	180.1	0.5776
8.4	28	m	Germany	52	75.42	170.0	0.6310

With these data we want to answer the following questions:

- Which of the customers' characteristics are important to be able to determine the consumption?
- How does the effect of important features look like?
- How does the model behave on unseen data?

What can we do to answer all the questions?


- Fit a linear model?
- Fit a linear model on each feature and build the ensemble?
- Fit a regularized linear model?
- Train a random forest?

What can we do to answer all the questions?

- Fit a linear model?
→ Not possible $p > n$
- Fit a linear model on each feature and build the ensemble?
→ Possible, but how should we determine important effects?
- Train a random forest?
→ Possible, but we want to interpret the effects.
- Fit a regularized linear model?
→ Possible, and basically the same as component-wise boosting.

Component-Wise Boosting

The Idea of Component-Wise Boosting



Why Component-Wise Boosting?

- Inherent (unbiased) feature selection.
- Resulting model is sparse since important effects are selected first and therefore it is able to learn in high-dimensional feature spaces ($p \gg n$).
- Parameters are updated iteratively. Therefore, the whole trace of how the model evolves is available.

About `compboost`

The `compboost` package is a fast and flexible framework for model-based boosting completely written in C++:

- With `mboost` as standard, we want to keep the modular principle of defining custom base-learner and losses.
- Completely written in C++ and exposed by `Rcpp` to obtain high performance and full memory control.
- R API is written in `R6` to provide convenient wrapper.
- Major parts of the `compboost` functionality are unit tested against `mboost` to ensure correctness.

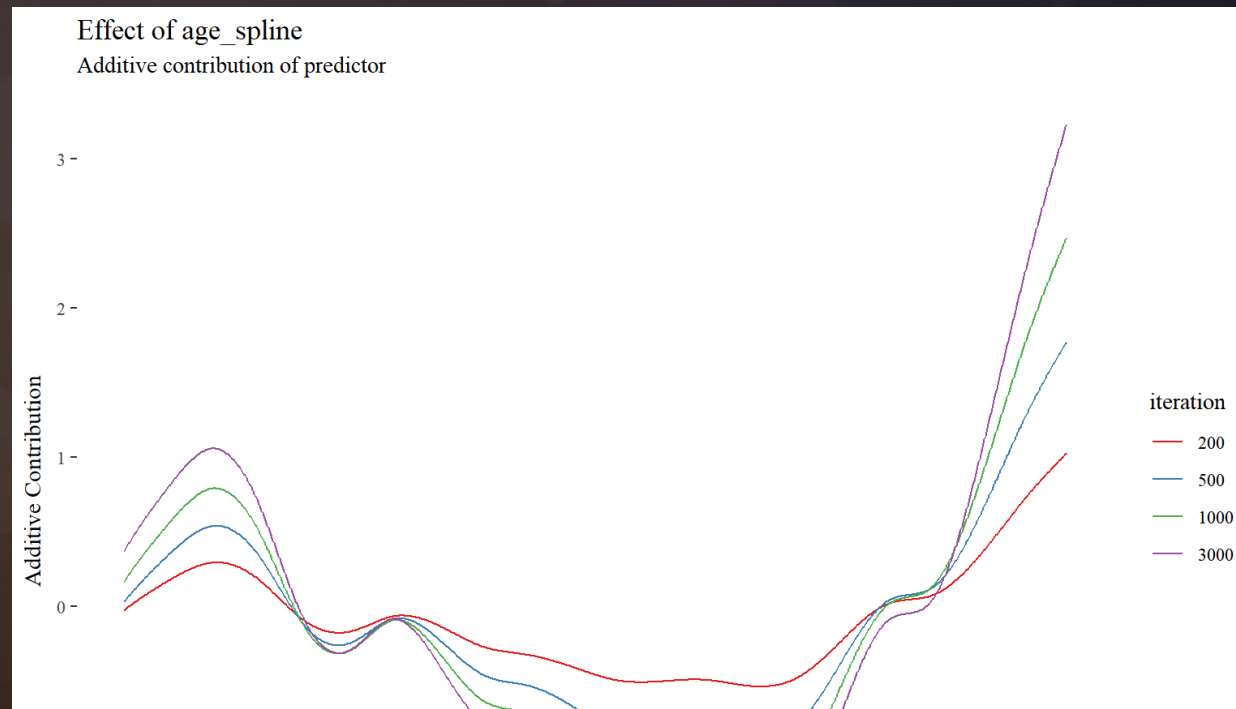
Applying compboost to the Use-Case

Quick Start With Wrapper Functions

```
cboost = boostSplines(data = mulled_wine_data[, -2], target = "mw_consumption",
  loss = LossQuadratic$new(), learning.rate = 0.005, iterations = 6000,
  trace = 600)
##      1/6000    risk = 5.6
##    600/6000    risk = 2.5
##   1200/6000    risk = 1.3
##   1800/6000    risk = 0.8
##   2400/6000    risk = 0.51
##   3000/6000    risk = 0.34
##
##   3600/6000    risk = 0.24
##   4200/6000    risk = 0.17
##   4800/6000    risk = 0.12
##   5400/6000    risk = 0.087
##   6000/6000    risk = 0.064
##
##
## Train 6000 iterations in 9 Seconds.
## Final risk based on the train set: 0.064
```

Effect Visualization

```
cboost$plot("age_spline", iters = c(200, 500, 1000, 3000))
```



Inbag and OOB Behavior

To get an idea, how the model behaves on unseen data we use 75 % as training data and the other 25 % of the data to calculate the out of bag (OOB) risk:

```
n_data = nrow(mulled_wine_data)
idx_train = sample(x = seq_len(n_data), size = n_data * 0.75)
idx_test = setdiff(x = seq_len(n_data), idx_train)
```

Define Model and Base-Learner the “Object-Oriented Style”

```
cboost = Comboost$new(data = mulled_wine_data[idx_train,-2],
  target = "mw_consumption", loss = LossQuadratic$new(),
  learning.rate = 0.005)

for (feature_name in setdiff(names(mulled_wine_data[, -2]), target)) {
  if (feature_name %in% c("gender", "country")) {

    cboost$addBaselearner(feature = feature_name, id = "category",
      bl.factory = BaselearnerPolynomial, intercept = FALSE)
  } else {
    cboost$addBaselearner(feature = feature_name, id = "spline",
      bl.factory = BaselearnerPSpline, degree = 3, n.knots = 10)
  }
}
```


OOB Data

To track the OOB risk we have to prepare the new data so that `compboost` knows the new data sources:

```
oob_data = cboost$prepareData(mulled_wine_data[idx_test,])  
oob_response = mulled_wine_data$mw_consumption[idx_test]
```

Define Logger

```
cboost$addLogger(logger = LoggerOobRisk, logger.id = "oob_risk",
  used.loss = LossQuadratic$new(), eps.for.break = 0,
  oob.data = oob_data, oob.response = oob_response)

cboost$addLogger(logger = LoggerTime, logger.id = "microseconds",
  max.time = 0, time.unit = "microseconds")

cboost$train(6000, trace = 1500)

##      1/6000      risk = 5.6  microseconds = 1    oob_risk = 5.7
## 1500/6000      risk = 0.99  microseconds = 2518925    oob_risk = 1.5
## 3000/6000      risk = 0.37  microseconds = 4848996    oob_risk = 1.1
## 4500/6000      risk = 0.17  microseconds = 7171435    oob_risk = 1.1
## 6000/6000      risk = 0.092  microseconds = 9495656    oob_risk = 1.2
##
##
## Train 6000 iterations in 9 Seconds.
## Final risk based on the train set: 0.092
```

Extract Inbag and OOB Data


```
inbag_trace = cboost$getInbagRisk()
head(inbag_trace)
## [1] 5.659 5.647 5.635 5.623 5.612 5.600

logger_data = cboost$getLoggerData()
head(logger_data)

##      _iterations microseconds oob_risk
## 1              1           1      5.671
## 2              2          1728      5.660
## 3              3          3374      5.649
## 4              4          5043      5.639
## 5              5          6697      5.628
## 6              6          8428      5.618
```

Visualization of Inbag and OOB Risk

```
oob_trace = logger_data[["oob_risk"]]\n\nrisk_data = data.frame(\n  risk = c(inbag_trace, oob_trace),\n  type = rep(c("inbag", "oob"), times = c(length(inbag_trace),\n    length(oob_trace))),\n  iter = c(seq_along(inbag_trace), seq_along(oob_trace))\n)\nggplot(risk_data, aes(x = iter, y = risk, color = type)) + geom_line()
```




Set Model to a Specific Iteration

```
cboost$train(2200)
cboost
## Component-Wise Gradient Boosting
##
## Trained on mulled_wine_data[idx_train, -2] with target mw_consumption
## Number of base-learners: 210
## Learning rate: 0.005
##
## Iterations: 2200
## Offset: 6.65
##
## LossQuadratic Loss:
##
## Loss function:  $L(y, x) = 0.5 * (y - f(x))^2$ 
##
##
```

Feature Importance

```
cboost$plotFeatureImportance(num.feats = 15L)
```



Using Custom Losses - Defining a Loss

Now we recognize that the customers do not buy the wine in liters but per cup, therefore it might be better to use a Poisson loss to take the counting data into account. For that reason we define a new custom loss:

```
lossPoi = function (truth, pred) {  
  return (-log(exp(pred)^truth * exp(-exp(pred)) / gamma(truth + 1)))  
}  
  
gradPoi = function (truth, pred) {  
  return (exp(pred) - truth)  
}  
  
constInitPoi = function (truth) {  
  return (log(mean.default(truth)))  
}  
  
# Define custom loss:  
my_custom_loss = LossCustom$new(lossPoi, gradPoi, constInitPoi)
```

Using Custom Losses - Using the Loss

```
cboost = boostSplines(data = mulled_wine_data[, -1],
  target = "mw_consumption_cups", loss = my_custom_loss,
  optimizer = OptimizerCoordinateDescent$new(),
  learning.rate = 0.005, iterations = 500, trace = 100,
  n.knots = 10, degree = 3)
##    1/500    risk = 5.3
## 100/500    risk = 3

## 200/500    risk = 2.7
## 300/500    risk = 2.6
## 400/500    risk = 2.6
## 500/500    risk = 2.5
##
##
## Train 500 iterations in 0 Seconds.
## Final risk based on the train set: 2.5
```


Further Functionalities

- Each logger can also be used as stopper, therefore we can use them for early stopping
- In combination with the custom loss, we can use the OOB logger to track performance measures like the AUC (in binary classification)
- Losses and base-learner can also be directly extended using C++ (see `getCustomCppExample()`)

From C++ to R

Rcpp

- Automated conversion between \mathbb{R} and C++ data structures, such as vectors, matrices, or even whole classes
- Seamless integration of **Armadillo** for linear algebra
- Complicated stuff like compilation, or again, the conversion between \mathbb{R} and C++ are handled automatically

C++ to R Wrap Up of compboost

R

```
> cboost = Compboost$new(...)
```

```
> cboost$addBaselearner(feats)
```

```
> cboost$addLogger(logger)
```

C++

```
> bl_list = BaselearnerFactoryList(...)  
> l_list = LoggerList(...)  
> loss = new LossQuadratic(...)  
> optim = new OptimizerCoordinateDesc(...)
```

```
> source = InMemoryData(feats)  
> target = InMemoryData()  
> bl_feats = BaselearnerPSpline(source, target)  
> bl_list.registerBaselearner(bl_feats)
```

```
> logger = LoggerOBRisk(...)  
> l_list.registerLogger(logger)
```

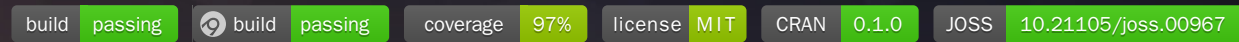
Challenges When Using C++ and Rcpp

- Saving object is not possible at the moment
- Memory managing is not easy \Rightarrow Segmentation faults or memory leaks may happen
- Exported API of classes is not very informative
- Debugging of C++ from R can be very annoying and time-consuming

What's Next?

- Better selection process of base-learner
- Speed up the training by parallel computations
- Greater functionality:
 - Functional data structures and loss functions
 - Unbiased feature selection
 - Effect decomposition into constant, linear, and non-linear

Thanks for your attention!



- Actively developed on GitHub:

<https://github.com/schalkdaniel/compboost>

<https://compboost.org/>

- Project page:

Credits

Slides were created with:

- **revealjs**
- **Font-Awesome:**
- **rmarkdown**
- **revealjs (R Package)**
- **Google Fonts**